



# K-Nearest Neighbors (KNN) on GPUs

Alexander Ingare

# Datasets

Name	Characteristics
MNIST	<ul style="list-style-type: none"><li>• Database of handwritten digits: 60k training, 10k testing</li><li>• 10 classes, 28x28 pixel, anti-aliased (grayscaled) images</li></ul>
CIFAR-10	<ul style="list-style-type: none"><li>• Subset if the 80 Million Tiny Images dataset</li><li>• 50k training images, 10k testing images</li><li>• 10 classes, 32x32 pixel, RGB images</li></ul>
STL-10	<ul style="list-style-type: none"><li>• Inspired by CIFAR, but used more for unsupervised learning</li><li>• 5k training images, 8k test images, 100k unlabeled images</li><li>• 10 classes, 96x96 pixel, RGB images</li></ul>




# GPU Platforms

Name	Characteristics
NVIDIA Tesla P100	<ul style="list-style-type: none"><li>• Oldest and slowest architecture</li><li>• Max (non-configurable) shared memory per SM (64 KB)</li></ul>
NVIDIA Tesla V100	<ul style="list-style-type: none"><li>• Newer architecture</li><li>• Max dynamic shared memory size per SM (96 KB)</li></ul>
NVIDIA A100	<ul style="list-style-type: none"><li>• Newest and fastest architecture</li><li>• Max dynamic shared memory size per SM (164 KB)</li></ul>



# K-Nearest Neighbors - Overview & Features

- Non-parametric supervised learning method most often used for classification with the output being class membership
- Objects are classified by a plurality vote of its neighbors
- Closest neighbor classification is determined by Euclidean-distance calculations
- Training examples are vectors in a multidimensional feature space each with a class label
- The training phase for KNN is simply loading the dataset training data (feature vectors and class labels of the training samples)





# High-Level Algorithm

1. Load and normalize image data into a vector with pixel values normalized within [0, 1]
2. Compare each test image to all training images
  - a. Compute the Euclidean distances between the two images
3. Identify the K Nearest Neighbors
  - a. Sort all training images by their distance to the test image (smallest to largest)
  - b. Select the top K closest training images
4. Do a majority vote on the neighbor labels
  - a. Return the label which appears most frequently amongst the nearest neighbors

$$d(x, y) = \sqrt{\sum_{i=1}^n (y_i - x_i)^2}$$

# Optimizations & Metrics

## Optimizations

Shared memory

GPU optimized sorting (thrust)

H2D image copy batching

## Metrics

Total & GPU Execution Time

Memory Allocation

KNN Accuracy

# MNIST Results

K=5	P100	V100	A100
Total Execution Time (s)	48.2293	9.1296	5.7931
GPU Execution Time (s)	47.7577	8.4668	4.6876
Memory Usage (MB)	182	182	182
Accuracy (%)	96.88	96.88	96.88

# CIFAR-10 Results

K=5	P100	V100	A100
Total Execution Time (s)	158.4035	38.3033	15.5068
GPU Execution Time (s)	157.2711	35.5580	13.1777
Memory Usage (MB)	588	588	588
Accuracy (%)	33.98	33.98	33.98

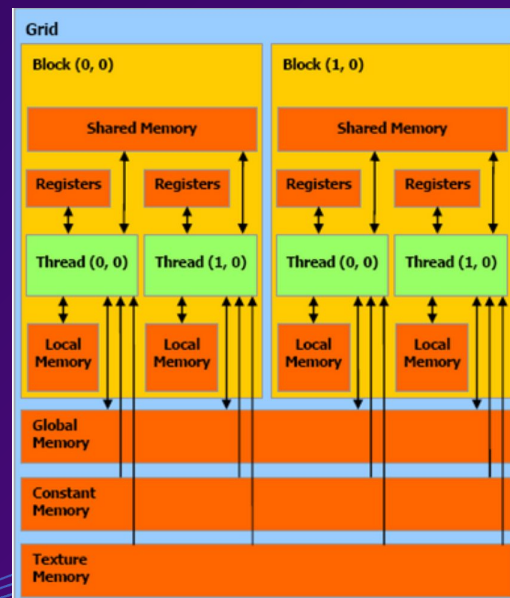


# STL-10 Results

K=5	P100	V100	A100
Total Execution Time (s)	62.2983	43.6378	44.2068
GPU Execution Time (s)	61.3881	42.6537	43.2905
Memory Usage (MB)	530	530	530
Accuracy (%)	26.925	26.925	26.925

# Setbacks

- Using shared memory for each test image, but test images are loaded linearly
  - No performance benefit since no computational reuse
  - Only a benefit if data is reused ***within the same block***
  - Only A100 GPUs can load a full STL-10 image! Tiling is needed
- Memory batching is overshadowed by sequential kernel executions



# Future

- Test multi-GPU program execution (if SLURM magically has them available)
- Use shared memory tiling for training images instead of test images
- Use a different image batching implementation
- Compare results for varying K values
- Compare unoptimized results to fully (working) optimized results
- Get more fine grained memory metrics (bandwidth, cache usage, etc...)
- Explore FP16 (instead of FP32) to utilize V100 & A100 Tensor Cores

The background is a solid dark purple. It features several abstract, flowing shapes in shades of blue and light purple. These shapes are composed of many thin, parallel lines that create a sense of movement and depth. Some areas have a grid-like pattern of intersecting lines, while others are more fluid and wavy. The overall effect is modern and dynamic.

# Q&A